


CS 344

LECTURE 7: GRAPHS


This is a graph:




A graph is a pair $\langle V, E \rangle$:

- V : vertices
- E : edges, where $E \subseteq V \times V$

This is a graph:





- $V = \{1, 2, 3\}$
- $E = \{(1, 2), (1, 3), (2, 3)\}$

GRAPH REPRESENTATION

- adjacency matrix
- adjacency list


ADJACENCY MATRIX

If we have $|V| = n$ vertices,

the adjacency matrix is an $n \times n$ matrix:

$$a_{ij} = \begin{cases} 1 & \text{if there is an edge from } v_i \text{ to } v_j \\ 0 & \text{otherwise} \end{cases}$$

ADJACENCY MATRIX



$$\begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

ADJACENCY MATRIX


- Can check for an edge in $O(1)$ time
- Takes $O(n^2)$ space!

ADJACENCY LIST

Each vertex has a linked list of its (outgoing) neighbors

That is, for each vertex u , it has a list of every v for
which $(u, v) \in E$

ADJACENCY LIST



Vertex	List
1	2, 3
2	1, 3
3	1, 2, 4
4	3

ADJACENCY LIST

- Have to walk a list to check for an edge
- Takes $O(|E|)$ space

SPARSITY

For a (connected) graph, we have that (roughly)

$$|V| \leq |E| \leq |V|^2$$

SPARSITY

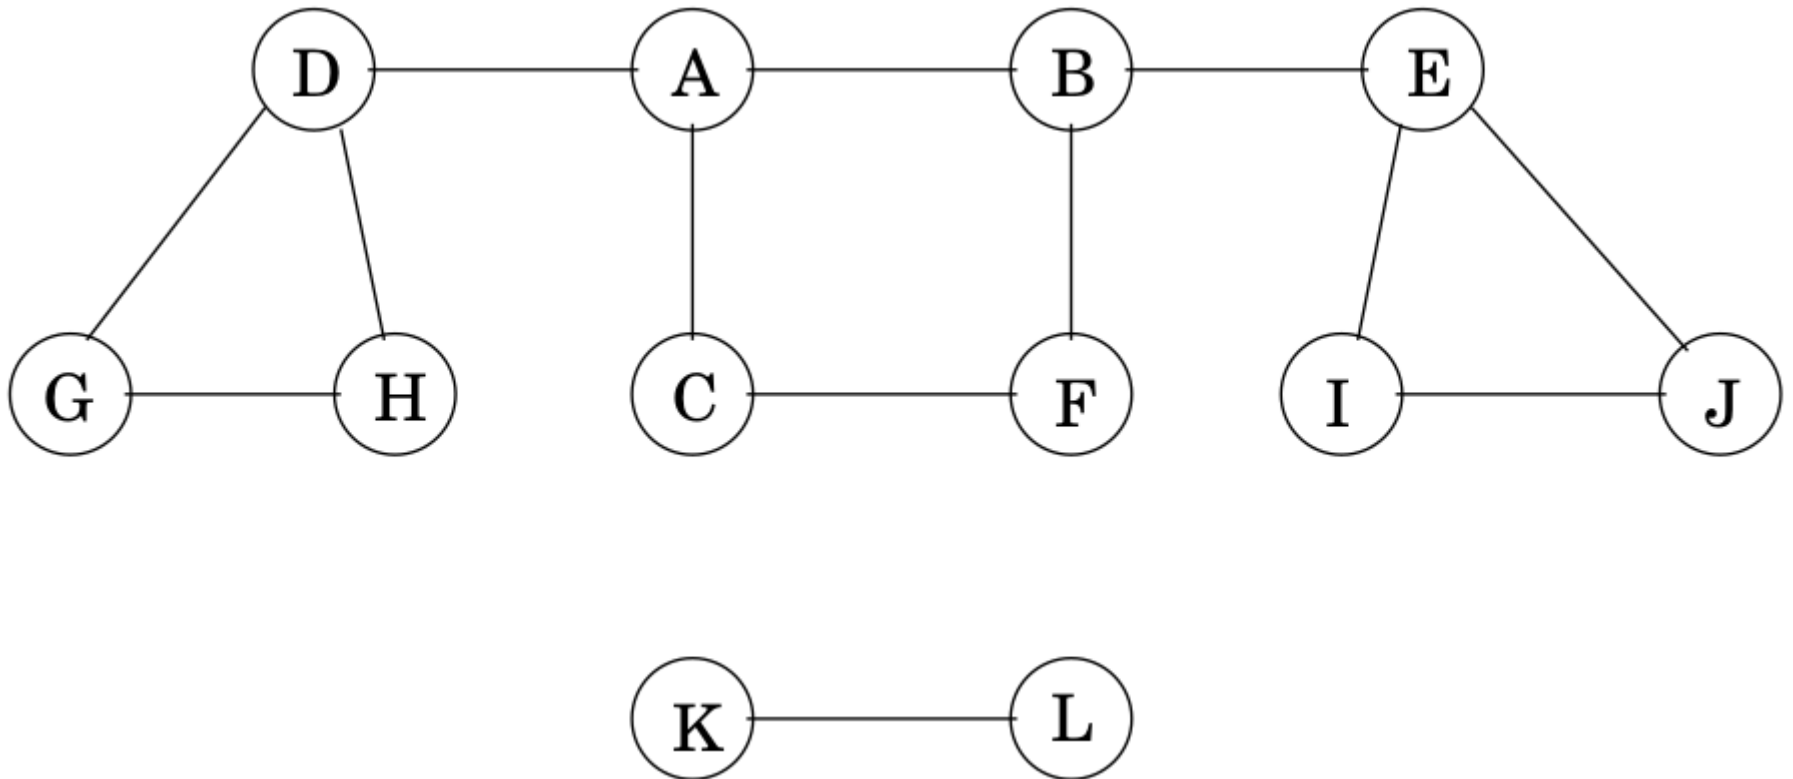
- Sparse: $|E|$ closer to $|V|$
- Dense: $|E|$ closer to $|V|^2$

SPARSITY


Which representation is better depends on the sparsity of the graph.

- For example, take the world-wide web
- Vertices are pages, edges are links
- $|V|$ is in the billions
- But most pages only link to a few other pages!

Suppose we have this graph:



Consider this as a maze:



To explore a maze we need two things:

- Chalk: to mark junctions you've visited
- String: to find your way back

To explore a graph we need two things:

- A variable to mark nodes you've visited
- A stack to find your way back

Let's find all nodes reachable from a particular node:

procedure explore(G, v)

Input: $G = (V, E)$ is a graph; $v \in V$

Output: $\text{visited}(u)$ is set to true for all nodes u reachable from v

$\text{visited}(v) = \text{true}$



$\text{previsit}(v)$

for each edge $(v, u) \in E$:


 if not $\text{visited}(u)$: $\text{explore}(u)$

$\text{postvisit}(v)$

Let's run `explore(A)` on the graph on the left:




Then we have two kinds of edges:



- tree edges: black lines
- back edges: dotted lines

Is the explore function correct?
Does it find all vertices reachable from v ?

Suppose not, that it misses vertex u :



More generally, correctness here means for any $k \geq 0$,
all nodes within k hops of v will be visited.

- Base case: $k = 0$
- Inductive step: If all nodes k hops away are visited, then all nodes $k + 1$ hops away are too

DEPTH-FIRST SEARCH

The explore function only visits nodes reachable from the starting point

To examine the rest of the graph, we can repeatedly call explore

procedure dfs (G)

for all $v \in V$:

 visited(v) = false

for all $v \in V$:

 if not visited(v): explore(v)

TIME TO RUN DFS

- $O(1)$ to mark node visited, call pre/postvisit
- Then loop through and scan adjacent edges





TIME TO RUN DFS

For all vertices together,

- $O(|V|)$ to mark nodes visited, call pre/postvisit
- Each edge (u, v) will be visited twice
 - once in $\text{explore}(u)$
 - once in $\text{explore}(v)$
- Therefore $O(|E|)$ work to scan edges



Total: $O(|V| + |E|)$

Running DFS on this graph (left) generates this forest (right):




An undirected graph is **connected** if there is a path from any vertex to any other.

In a disconnected graph, each connected subgraph is called a **connected component**.




Connected components:




Connected components:

- $\{A, B, E, I, J\}$



Connected components:

- $\{A, B, E, I, J\}$
- $\{C, D, G, H, K, L\}$



Connected components:

- $\{A, B, E, I, J\}$
- $\{C, D, G, H, K, L\}$
- $\{F\}$

procedure previsit(v)
 $ccnum[v] = cc$





- Initialize to 0
- Increment each time DFS calls explore

Let's add a counter to see when we enter/leave nodes:


```
procedure previsit( $v$ )  
pre[ $v$ ] = clock  
clock = clock + 1
```

```
procedure postvisit( $v$ )  
post[ $v$ ] = clock  
clock = clock + 1
```



Then our forest now looks like this:




DIRECTED GRAPHS



DFS ON DIRECTED GRAPHS




Terminology:




- ***A* is the root**
- ***E* has descendants *F*, *G*, and *H***
- ***E* is an ancestor of *F*, *G*, and *H***
- ***C* is the parent of *D***
- ***D* is the child of *C***

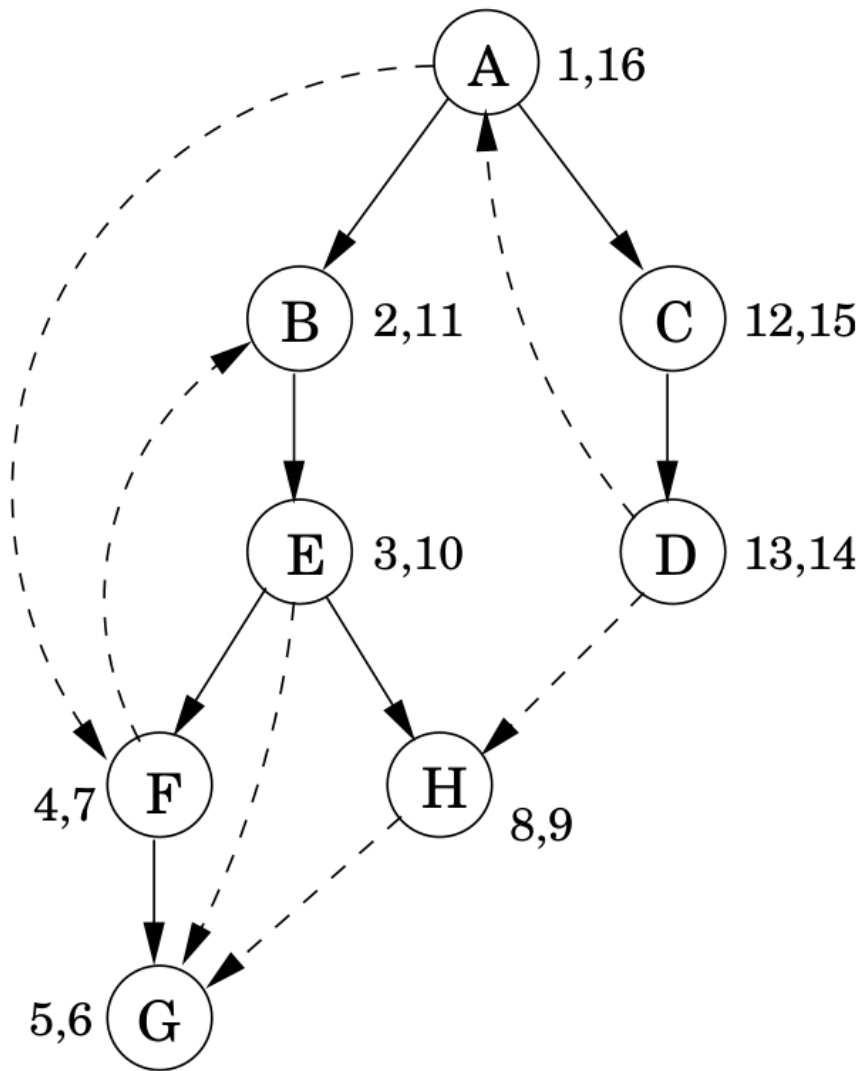
We can also have finer-grained distinctions on edges in the generated tree:



- tree edges
- forward edges
- back edges
- cross edges



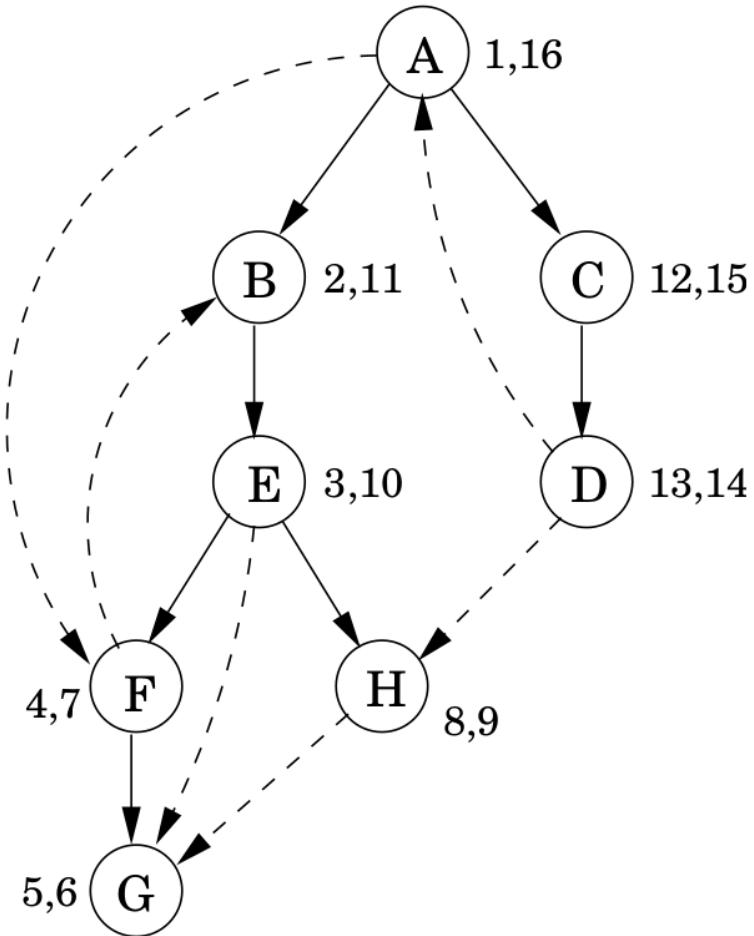
- tree edges: part of the DFS forest
- forward edges: node to nonchild descendent
- back edges: node to ancestor
- cross edges: lead to neither descendent nor ancestor



How many


- Forward edges?
- Back edges?
- Cross edges?

These relationships can be inferred from the pre and post numbers!



Vertex u is an ancestor of vertex v :





Edge categories:

pre/post <i>ordering for</i> (u, v)				<i>Edge type</i>
$\left[\begin{array}{c} \\ u \end{array} \right.$	$\left[\begin{array}{c} \\ v \end{array} \right.$	$\left. \begin{array}{c} \\ v \end{array} \right]$	$\left. \begin{array}{c} \\ u \end{array} \right]$	Tree/forward
$\left[\begin{array}{c} \\ v \end{array} \right.$	$\left[\begin{array}{c} \\ u \end{array} \right.$	$\left. \begin{array}{c} \\ u \end{array} \right]$	$\left. \begin{array}{c} \\ v \end{array} \right]$	Back
$\left[\begin{array}{c} \\ v \end{array} \right.$	$\left. \begin{array}{c} \\ v \end{array} \right]$	$\left[\begin{array}{c} \\ u \end{array} \right.$	$\left. \begin{array}{c} \\ u \end{array} \right]$	Cross

CYCLES

A cycle is a circular path $v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_k \rightarrow v_0$.

A graph without cycles is acyclic.

DFS AND CYCLES

A directed graph has a cycle if and only if DFS reveals a back edge.

A directed graph has a cycle if and only if DFS reveals a back edge.

A directed graph has a cycle if and only if DFS reveals a back edge.

- (\Leftarrow) If (u, v) is a back edge,

A directed graph has a cycle if and only if DFS reveals a back edge.

- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u

A directed graph has a cycle if and only if DFS reveals a back edge.

- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u
 - That path plus the back edge is a cycle

A directed graph has a cycle if and only if DFS reveals a back edge.

- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u
 - That path plus the back edge is a cycle
- (\Rightarrow) If $v_0 \rightarrow \dots \rightarrow v_k \rightarrow v_0$ is a cycle:

A directed graph has a cycle if and only if DFS reveals a back edge.

- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u
 - That path plus the back edge is a cycle
- (\Rightarrow) If $v_0 \rightarrow \dots \rightarrow v_k \rightarrow v_0$ is a cycle:
 - Consider the first node v_i to be discovered

A directed graph has a cycle if and only if DFS reveals a back edge.


- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u
 - That path plus the back edge is a cycle
- (\Rightarrow) If $v_0 \rightarrow \dots \rightarrow v_k \rightarrow v_0$ is a cycle:
 - Consider the first node v_i to be discovered
 - It will reach all other nodes on the cycle

A directed graph has a cycle if and only if DFS reveals a back edge.


- (\Leftarrow) If (u, v) is a back edge,
 - There is a path from v to u
 - That path plus the back edge is a cycle
- (\Rightarrow) If $v_0 \rightarrow \dots \rightarrow v_k \rightarrow v_0$ is a cycle:
 - Consider the first node v_i to be discovered
 - It will reach all other nodes on the cycle
 - The edge $v_{i-1} \rightarrow v_i$ will be a back edge

DIRECTED ACYCLIC GRAPH (DAG)


Dags are good for modeling hierarchies or dependencies (e.g., course prerequisites).



Given a dag, we may want to linearize (topologically sort) the nodes.



Given a dag, we may want to linearize (topologically sort) the nodes.





One possibility: B, A, D, C, E, F

How can we linearize a dag algorithmically?



How can we linearize a dag algorithmically?

List nodes in decreasing order of post numbers

List nodes in decreasing order of post numbers



List nodes in decreasing order of post numbers



B, D, A, C, F, E

Given a dag, we can define two kinds of special nodes:

- source: a node with no incoming edges
- sink: a node with no outgoing edges

Given a dag, we can define two kinds of special nodes:

- source: a node with no incoming edges
- sink: a node with no outgoing edges

Then a linearization must start with a source and end with a sink.

Another algorithm for linearization:

Another algorithm for linearization:

- find a source

Another algorithm for linearization:

- find a source
- output it

Another algorithm for linearization:

- find a source
- output it
- delete it

Another algorithm for linearization:


- find a source
- output it
- delete it
- repeat until graph is empty

STRONGLY-CONNECTED COMPONENTS


Two nodes u and v are connected if there is a u - v path as well as a v - u path.

We can partition a directed graph into a set **strongly-connected components** (SCCs), where all vertices are connected.

Dotted lines indicate the SCCs:



Then we can shrink each component to a single meta-vertex:



A directed graph is a dag of its strongly-connected components.

How can we decompose a graph to SCCs?

- If we found a node in a sink SCC
 - explore would find all nodes in its SCC
- Then we could remove it and repeat

How can we decompose a graph to SCCs?
But it's easier to find a node in a source SCC:

How can we decompose a graph to SCCs?
But it's easier to find a node in a source SCC:
node with the highest post number in DFS

More generally,

- if C and C' are SCCs,
- and there's a $C \rightarrow C'$ edge,
- then the highest post number in $C >$ highest in C'

This means SCCs can be linearized by decreasing
highest post numbers.

(a generalization of linearization of dags)

How can we decompose a graph to SCCs?

We can find a source SCC, but we need a sink SCC

How can we decompose a graph to SCCs?


We can find a source SCC, but we need a sink SCC

Transform the graph G into the reverse graph G^R !


How can we decompose a graph to SCCs?

- Compute G^R
- Run DFS on G^R
- Find the connected components using explore, in decreasing order of post numbers


Suppose we have this graph:



Then G^R is:



And a DFS gives these post numbers:



Then we can separate this into SCCs $\{C, D\}$ and $\{A, B\}$:

