


CS 344



LECTURE 8

GRAPHS, PATHS, AND BFS


Suppose we have a graph, and do DFS search from S :



We should get the following tree:




We should get the following tree:




But if we were interested in getting to C , this gives a rather inefficient route of $S \rightarrow A \rightarrow B \rightarrow C$ instead of $S \rightarrow C$

Imagine the graph as a physical set of marbles connected by string:



When we pick up the graph by S , we see immediately how to get to C :



BREADTH-FIRST SEARCH (BFS)

- Proceed layer by layer
- Find layer $d + 1$ by scanning neighbors of layer d

procedure `bfs`(G, s)

Input: Graph $G = (V, E)$, directed or undirected; vertex $s \in V$

Output: For all vertices u reachable from s , `dist`(u) is set to the distance from s to u .

for all $u \in V$:

`dist`(u) = ∞

`dist`(s) = 0

$Q = [s]$ (queue containing just s)

while Q is not empty:

$u = \text{eject}(Q)$


 for all edges $(u, v) \in E$:

 if `dist`(v) = ∞ :

`inject`(Q, v)


`dist`(v) = `dist`(u) + 1

Let's run BFS on our graph, starting at S :



Let's run BFS on our graph, starting at S :

Order of visitation	Queue contents after processing node
S	$[S]$
A	$[A\ C\ D\ E]$
C	$[C\ D\ E\ B]$
D	$[D\ E\ B]$
E	$[E\ B]$
B	$[B]$
	$[\]$



IS BFS CORRECT?

That is, for each d , at some point:

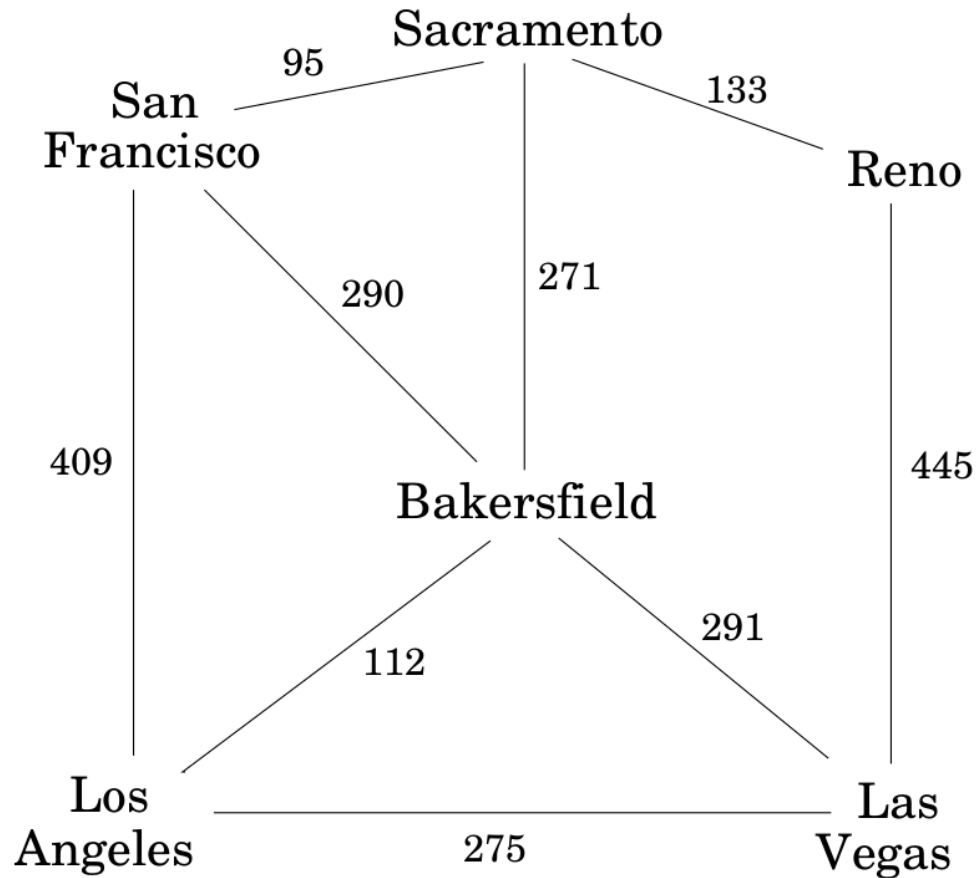
- all nodes with distance $\leq d$ are correctly set
- all other nodes have distance set to ∞
- the queue contains exactly the nodes at distance d

BFS RUNNING TIME

- Each vertex put onto queue once
- Examine each edge once (directed) or twice (undirected)

$$O(|V| + |E|)$$


Edges are often weighted with some kind of length:




We could compute shortest paths using BFS:

- Break edges into unit lengths with dummy nodes!

Suppose this is our graph:




We can break up each edge based on its length:




And then run BFS as usual

But what if this was our graph?




Then BFS is a slow, boring process:

G' :



Then BFS is a slow, boring process:


G' :



What if we could tell BFS to wake us up when it gets to the interesting part?


Wake us up when it gets to the interesting part:

G :



Wake us up when it gets to the interesting part:


G :



- ETAs for A and B are 100 and 200

Wake us up when it gets to the interesting part:


G :



- ETAs for A and B are 100 and 200
- Set alarms for 100 and 200

Wake us up when it gets to the interesting part:


G :



- ETAs for A and B are 100 and 200
- Set alarms for 100 and 200
- Wake up at time 100

Wake us up when it gets to the interesting part:


G :



- ETAs for A and B are 100 and 200
- Set alarms for 100 and 200
- Wake up at time 100
- New ETA for B is 150

Wake us up when it gets to the interesting part:

G :



- ETAs for A and B are 100 and 200
- Set alarms for 100 and 200
- Wake up at time 100
- New ETA for B is 150
- Change B 's alarm to 150

Wake us up when it gets to the interesting part:

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Say the next alarm goes off at time T for node u

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Say the next alarm goes off at time T for node u

- Then the distance from s to u is T

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Say the next alarm goes off at time T for node u

- Then the distance from s to u is T
- For each neighbor v of u :

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Say the next alarm goes off at time T for node u

- Then the distance from s to u is T
- For each neighbor v of u :
 - If there's no alarm for v , set one for $T + l(u, v)$

Wake us up when it gets to the interesting part:

- Set alarm for node s at time 0
- Repeat until no more alarms:

Say the next alarm goes off at time T for node u

- Then the distance from s to u is T
- For each neighbor v of u :
 - If there's no alarm for v , set one for $T + l(u, v)$
 - If v 's alarm is later than that, reduce it to this

This is essentially Dijkstra's algorithm!

(This is essentially Dijkstra)



DIJKSTRA'S ALGORITHM

Given a graph G and a starting vertex s ,
find shortest paths to all reachable vertices

DIJKSTRA'S ALGORITHM

We need to use a **priority queue** with these operations:

- insert
- decrease-key
- delete-min
- make-queue

procedure `dijkstra`(G, l, s)

Input: Graph $G = (V, E)$, directed or undirected;
 positive edge lengths $\{l_e : e \in E\}$; vertex $s \in V$

Output: For all vertices u reachable from s , `dist`(u) is set
 to the distance from s to u .

for all $u \in V$:

`dist`(u) = ∞

`prev`(u) = nil

`dist`(s) = 0

$H = \text{makequeue}(V)$ (using `dist`-values as keys)

while H is not empty:

$u = \text{deletemin}(H)$


 for all edges $(u, v) \in E$:

 if `dist`(v) > `dist`(u) + $l(u, v)$:


`dist`(v) = `dist`(u) + $l(u, v)$

`prev`(v) = u


`decreasekey`(H, v)




A: 0	D: ∞
B: 4	E: ∞
C: 2	



A: 0	D: 6
B: 3	E: 7
C: 2	




A: 0	D: 5
B: 3	E: 6
C: 2	



A: 0	D: 5
B: 3	E: 6
C: 2	


Finally, we get this tree for paths from A :



Dijkstra's algorithm is basically just BFS:


- Instead of a regular queue,
- use a priority queue to account for lengths

Another way of looking at this:




- Start at s
- Grow the "known region" R to larger distances
- Next node added to R is the one closest to s

How do we find the next node v outside R ?



- Consider shortest path $s \rightarrow \dots \rightarrow u \rightarrow v$
- u closer to s than v
- $u \in R$
- So $s \rightarrow v$ path extends a currently known path by one edge

How do we find the next node v outside R ?



- Try all single-edge extensions, find the shortest path
- Its endpoint is v

That is, v is the node outside R where $\text{distance}(s, u) + l(u, v)$ is minimized, for all $u \in R$

```
Initialize  $\text{dist}(s)$  to 0, other  $\text{dist}(\cdot)$  values to  $\infty$   
 $R = \{ \}$  (the ''known region'')  
while  $R \neq V$ :  
    Pick the node  $v \notin R$  with smallest  $\text{dist}(\cdot)$   
    Add  $v$  to  $R$   
    for all edges  $(v, z) \in E$ :  
        if  $\text{dist}(z) > \text{dist}(v) + l(v, z)$ :  
             $\text{dist}(z) = \text{dist}(v) + l(v, z)$ 
```


DIJKSTRA'S RUNNING TIME

- makequeue: at most $|V|$ insert operations
- $|V|$ deletemin operations
- $|V| + |E|$ insert/decreasekey operations

Time depends on implementation, but if we use a
binary heap:

$$O((|V| + |E|) \log |V|)$$

How can we implement the priority queue?

- Array
- Binary heap
- d -ary heap
- Fibonacci heap

ARRAY

- insert, decreasekey: $O(1)$
- deletemin: $O(n)$


BINARY HEAP

- Complete binary tree (except last level)
- Key value \leq that of its children


BINARY HEAP

- insert:
 - place node at bottom of tree and "bubble up"
- decreasekey:
 - "bubble up" (it's already in the tree)
- deletemin:
 - remove (and return) root
 - move last element to root
 - "sift down"


Let's say we have this initial heap:




We insert 7:




Bubble up:




Bubble up:




Now let's run delete-min:




Move last node to root:



Sift down:



Sift down:



d -ARY HEAP

- Like binary heap, but with d children for each node
- Height reduces to


$$\Theta(\log_d n) = \Theta((\log n)/(\log d))$$

- Insert/decreasekey slightly faster
- Deletemin slightly slower



Which implementation is best?

Implementation	deletemin	insert/ decreasekey	$ V \times \text{deletemin} + (V + E) \times \text{insert}$
Array	$O(V)$	$O(1)$	$O(V ^2)$
Binary heap	$O(\log V)$	$O(\log V)$	$O((V + E) \log V)$
d -ary heap	$O(\frac{d \log V }{\log d})$	$O(\frac{\log V }{\log d})$	$O((V \cdot d + E) \frac{\log V }{\log d})$
Fibonacci heap	$O(\log V)$	$O(1)$ (amortized)	$O(V \log V + E)$

Suppose now that the graph has negative edges:



Can we just shift everything into the positive range?




Note that the distances in Dijkstra's algorithm are always \leq the true distance.

procedure update $((u, v) \in E)$
 $\text{dist}(v) = \min\{\text{dist}(v), \text{dist}(u) + l(u, v)\}$

- If u is the second-last node in shortest path to v , gives exact distance
- Cannot set $\text{dist}(v)$ smaller than the true distance

Dijkstra's algorithm can be seen as a sequence of these updates

Consider the shortest path from s to some node t :




- This path has at most $|V| - 1$ edges
- If we did updates in this order:

$$(s, u_1), (u_1, u_2), (u_2, u_3), \dots, (u_k, t)$$

we would get the correct distance to t .

Consider the shortest path from s to some node t :



- How can we update the right edges in the right order, without already knowing the shortest paths?
- Just update all edges $|V| - 1$ times!

\Rightarrow Bellman-Ford algorithm, $O(|V| \cdot |E|)$.

Bellman-Ford algorithm:

procedure shortest-paths (G, l, s)


Input: Directed graph $G = (V, E)$;
 edge lengths $\{l_e : e \in E\}$ with no negative cycles;
 vertex $s \in V$

Output: For all vertices u reachable from s , $\text{dist}(u)$ is set
 to the distance from s to u .

for all $u \in V$:
 $\text{dist}(u) = \infty$
 $\text{prev}(u) = \text{nil}$


$\text{dist}(s) = 0$
repeat $|V| - 1$ times:
 for all $e \in E$:
 update(e)

Bellman-Ford algorithm:



	Iteration							
Node	0	1	2	3	4	5	6	7
S	0	0	0	0	0	0	0	0
A	∞	10	10	5	5	5	5	5
B	∞	∞	∞	10	6	5	5	5
C	∞	∞	∞	∞	11	7	6	6
D	∞	∞	∞	∞	∞	14	10	9
E	∞	∞	12	8	7	7	7	7
F	∞	∞	9	9	9	9	9	9
G	∞	8	8	8	8	8	8	8

What if there is a negative cycle?



Two kinds of graphs that cannot have negative cycles:

- graphs with no negative edges
- dags

SHORTEST PATHS IN DAGS

Key idea:

- In any path, the vertices appear in increasing linearized order

SHORTEST PATHS IN DAGS

- Linearize the dag using DFS
- Visit the vertices in sorted order
- Update all outgoing edges

procedure dag-shortest-paths(G, l, s)

Input: Dag $G = (V, E)$;

 edge lengths $\{l_e : e \in E\}$; vertex $s \in V$

Output: For all vertices u reachable from s , $\text{dist}(u)$ is set
 to the distance from s to u .

for all $u \in V$:

$\text{dist}(u) = \infty$

$\text{prev}(u) = \text{nil}$

$\text{dist}(s) = 0$

Linearize G



for each $u \in V$, in linearized order:

 for all edges $(u, v) \in E$:

$\text{update}(u, v)$


Recap: How can we linearize a dag algorithmically?

- List nodes in decreasing order of post numbers




Order: B, D, A, C, F, E

Now let's add some lengths:




Order: B, D, A, C, F, E




B, D, A, C, F, E

	0
A	∞
B	0
C	∞
D	∞
E	∞
F	∞




B, D, A, C, F, E

	0	1
A	∞	3
B	0	0
C	∞	∞
D	∞	1
E	∞	∞
F	∞	∞




B, D, A, C, F, E

	0	1	2
A	∞	3	3
B	0	0	0
C	∞	∞	6
D	∞	1	1
E	∞	∞	∞
F	∞	∞	∞



B, D, A, C, F, E

	0	1	2	3
A	∞	3	3	3
B	0	0	0	0
C	∞	∞	6	5
D	∞	1	1	1
E	∞	∞	∞	∞
F	∞	∞	∞	∞



B, D, A, C, F, E

	0	1	2	3	4
A	∞	3	3	3	3
B	0	0	0	0	0
C	∞	∞	6	5	5
D	∞	1	1	1	1
E	∞	∞	∞	∞	7
F	∞	∞	∞	∞	6